Datlas Challenge

Team: Enigma

Ramon Díaz
Daniela Gómez
Michael Zenkl
Jorge Ayala

Objective

To obtain insights from Nuevo Leon’s car accidents with the use of data analysis and machine learning methods.

Methodology

- CRISP-DM Methodology (Cross-industry standard process for data mining)
  - Business Understanding
  - Data Understanding
  - Data Preparation
  - Modeling
  - Evaluation
  - Deploying

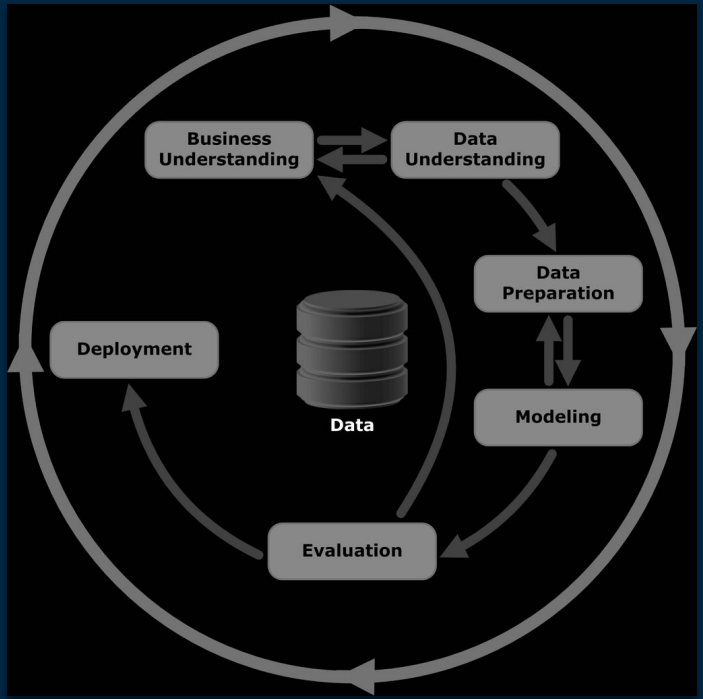

Data Cleaning

- Missing Values (statistical imputation)
- Binned Variables with Multiple Labels (Model Years)
- Removed Special Characters (Car Colors)

Feature Engineering

- Climate conditions
- Type of road identification
- Nearby pedestrian crossings
- Geohashing of coordinates in different areas (from 5 to 8 hash characters)
- Identification of holidays

Data Preprocessing

- Standardization : MaxMin method
- Recursive Feature Elimination for Feature Selection

Tools

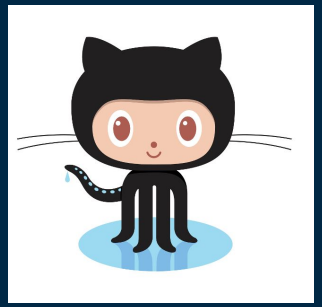

https://github.com/danisha20/project_DATLAS

Methods

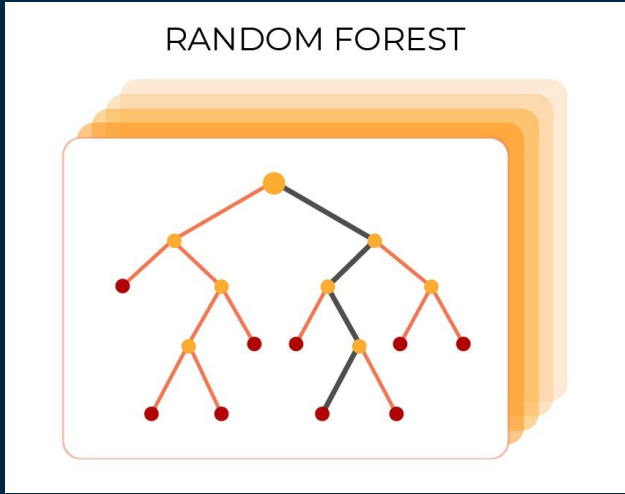

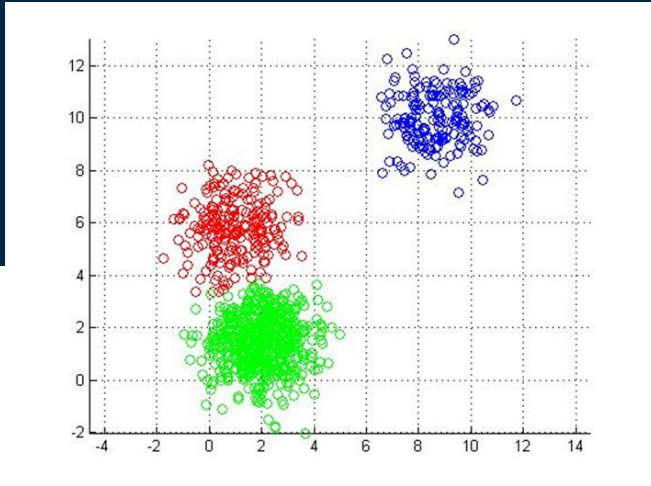

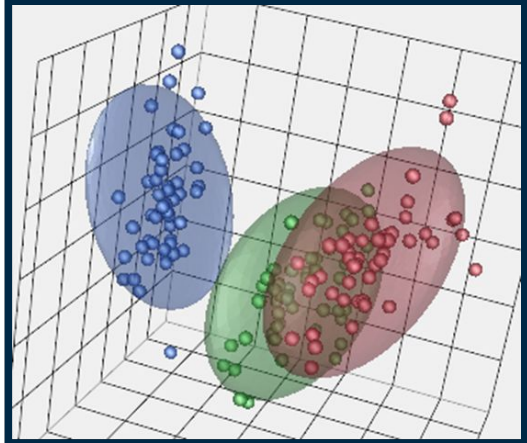

K-means
K-modes

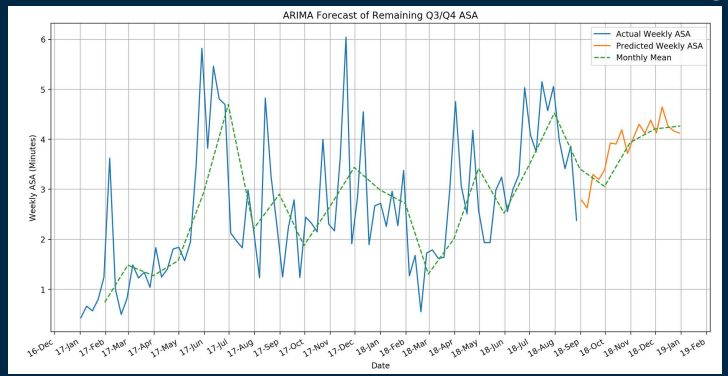

EDA

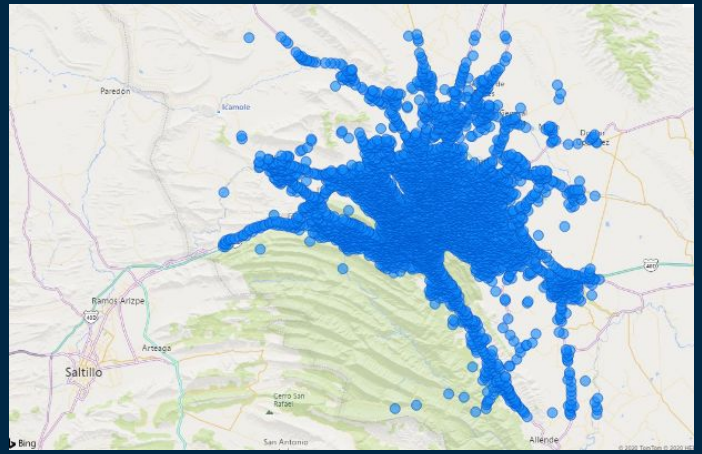
Parras
Ramos Arizpe
Arteaga
Saltillo
Cerro San Rafael
San Antonio
Allende
Bing

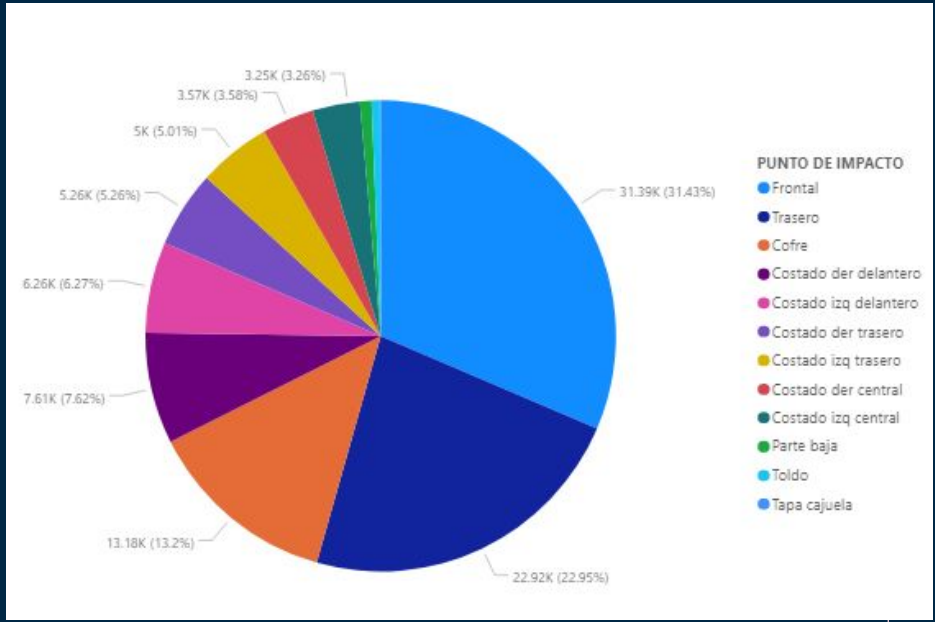
3.25K (3.26%)
3.57K (3.58%)
5K (5.01%)
5.26K (5.26%)
6.26K (6.27%)
7.61K (7.62%)
13.18K (13.2%)
22.92K (22.95%)
31.39K (31.43%)
PUNTO DE IMPACTO
Frontal
Trasero
Cofre
Costado der delantero
Costado izq delantero
Costado der trasero
Costado izq trasero
Costado der central
Costado izq central
Parte baja
Toldo
Tapa cajuela

EDA

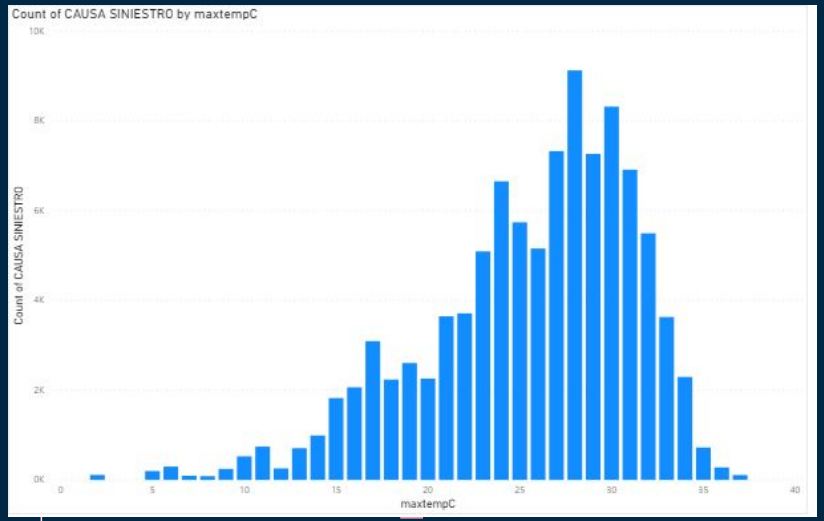

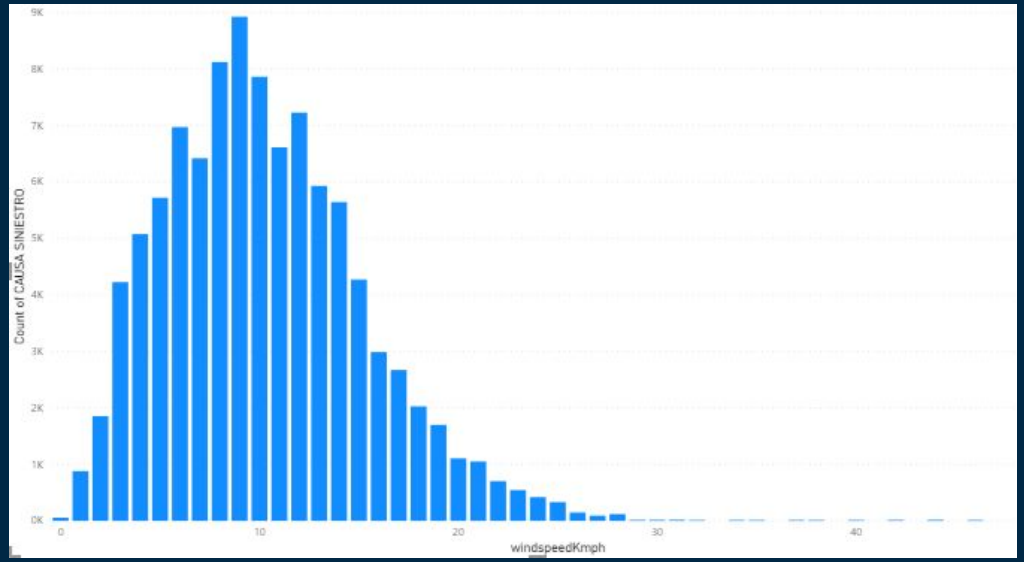

Mean max temp = 25.219°C
SD max temp = 5.745°C
Mean winds peed = 10.258 kmph
SD wind speed =5.029 kmph

EDA

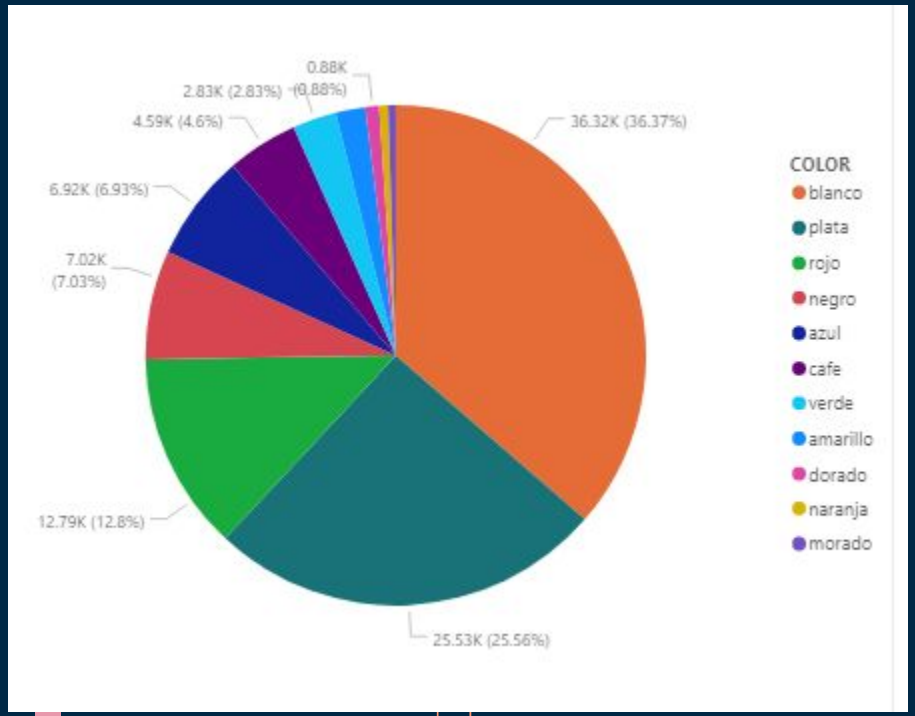
0.88K
(0.88%)
2.83K (2.83%)
4.59K (4.6%)
36.32K (36.37%)
6.92K (6.93%)
7.02K
(7.03%)
12.79K (12.8%)
25.53K (25.56%)
COLOR
blanco
plata
rojo
negro
azul
cafe
verde
amarillo
dorado
naranja
morado

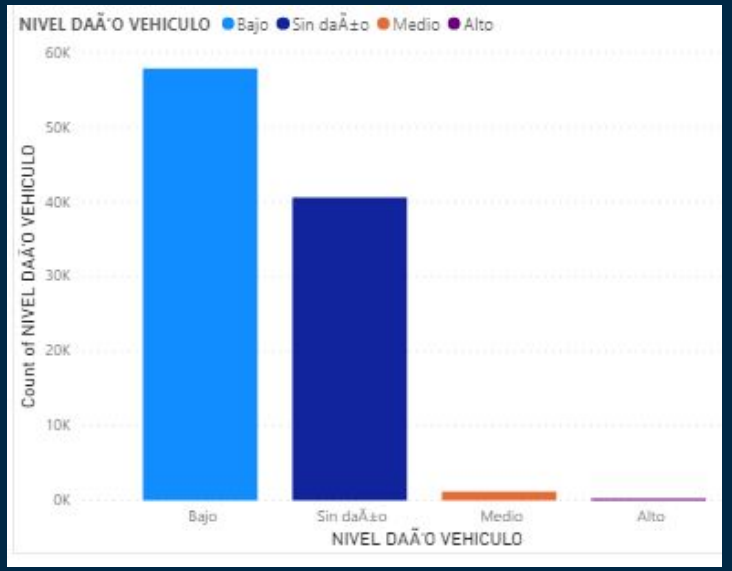
NIVEL DAÃ'O VEHICULO
Bajo
Sin daÃ±o
Medio
Alto
60K
50K
40K
30K
20K
10K
0K
Count of NIVEL DAÃ'O VEHICULO
Bajo
Sin daÃ±o
Medio
Alto
NIVEL DAÃ'O VEHICULO

EDA

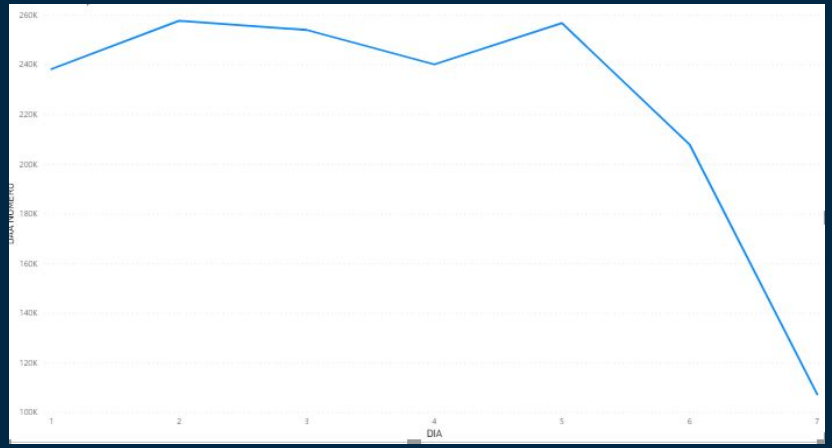
260K
240K
220K
200K
180K
160K
140K
120K
100K
1
2
3
4
5
6
7
DIA

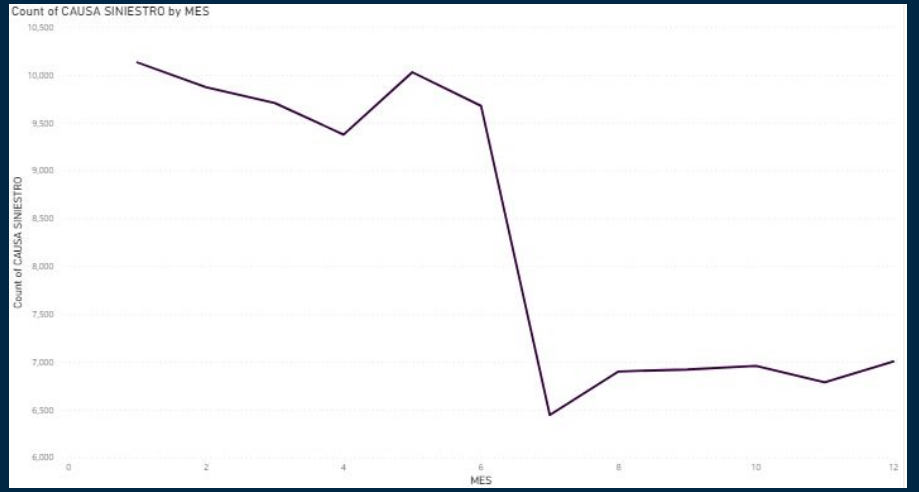
Count of CAUSA SINIESTRO by MES
10,500
10,000
9,500
9,000
8,500
8,000
7,500
7,000
6,500
6,000
Count of CAUSA SINIESTRO
0
2
4
6
8
10
12
MES

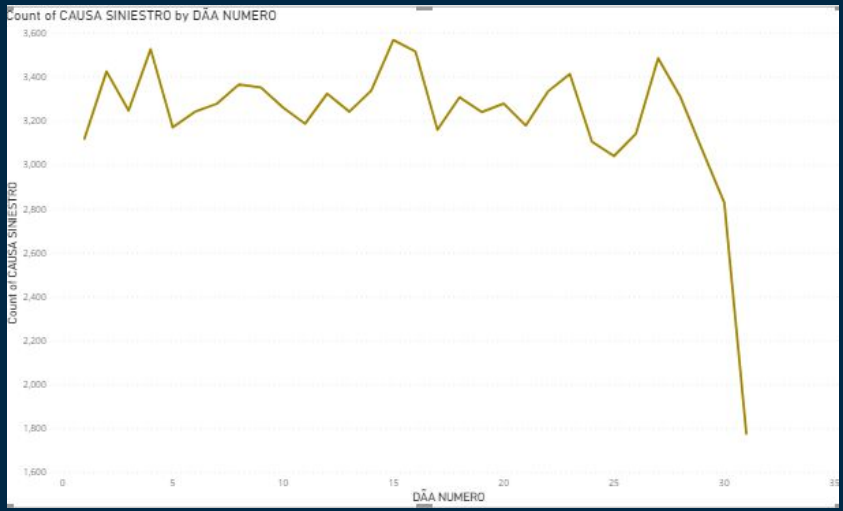
Count of CAUSA SINIESTRO by DÃA NUMERO
3,600
3,400
3,200
3,000
2,800
2,600
2,400
2,200
2,000
1,800
1,600
Count of CAUSA SINIESTRO
0
5
10
15
20
25
30
DÃA NUMERO

Machine Learning - Clustering

Clustering was made with kMeans and kModes (a variant for categorical datasets).
It considered:

- Location of accident
- Type of accident
- Car data
- Climate data
- Visibility
- If it was a holiday
- Type of road where it happened
- If it happened near a pedestrian crossing
- If it happened at night

Machine Learning - Classification

Validity Index using supervised Classifiers.

k-Means clustering was validated with 0.999 AUC with Random Forest classification.

The importance of the attributes in the construction of the random forest were as follows:

NIVEL DAÑO VEHICULO_2	Bajo
NIVEL DAÑO VEHICULO_1	Sin daño
TIPO VEHICULO_1	Camión
TIPO VEHICULO_2	Auto
MODEL_YEAR_5	older
NIVEL DAÑO VEHICULO_3	Medio
PUNTO DE IMPACTO_2	Frontal

Machine Learning - Classification

What does this mean?
The best way to segment the dataset is by accident severity.
Cars are more likely to get damaged than trucks.
Older cars are more likely to get damaged than newer ones.

This information could be useful for car insurance companies in determining how likely it is for the car to be damaged considering its make, it's age, and where it mostly drives.

Machine Learning – Logistic Regression

- Logistic regression not entirely suitable for this task
  - Relationship between variables not linear.
- Recursive Feature Elimination employed, but low accuracy achieved
- Adjusted R-Squared > 0.3

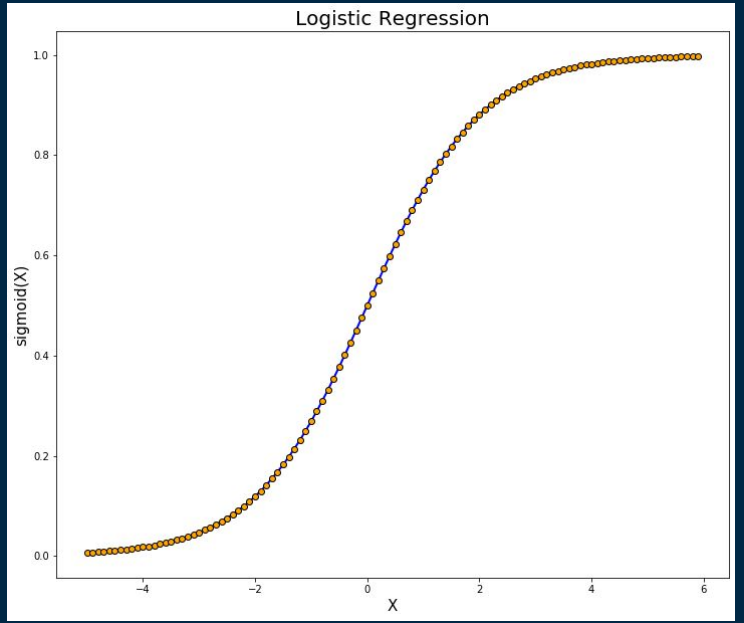

Machine Learning - Scikitlearn Random Forest

- Given the large amount of variables, Random Forest was suitable for the task
- 95.43% percent of accuracy
- Misprediction of ~5 accidents per day
- Geohash zones of 39.1km×19.5km (4 hash characters)

Machine Learning - Time-Series Analysis

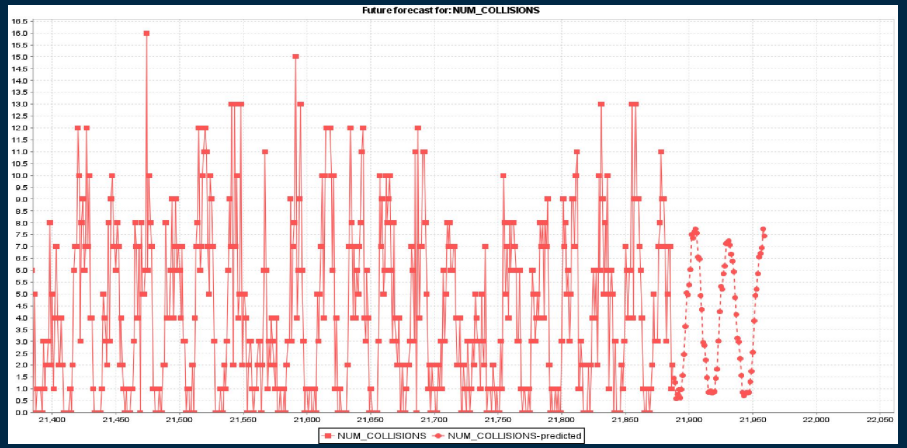

Random Forest with 72 hours ahead.

```
=== Evaluation on training data ===
```

Target	1-step-ahead	2-steps-ahead	3-steps-ahead	4-steps-ahead	5-steps-ahead	6-steps-ahead	7-steps-ahead	8-steps-ahead	9-steps-ahead
NUM_COLLISIONS									
N	21875	21874	21873	21872	21871	21870	21869	21868	21867
Mean absolute error	0.7783	1.1565	1.4615	1.6956	1.8803	2.0317	2.1678	2.2892	2.4009
Root mean squared error	1.0969	1.6435	2.0798	2.3838	2.6123	2.8054	2.9767	3.1533	3.3206

62-steps-ahead	63-steps-ahead	64-steps-ahead	65-steps-ahead	66-steps-ahead	67-steps-ahead	68-steps-ahead	69-steps-ahead	70-steps-ahead	71-steps-ahead	72-steps-ahead
21814	21813	21812	21811	21810	21809	21808	21807	21806	21805	21804
3.2399	3.2465	3.2412	3.2143	3.1891	3.1579	3.1483	3.1413	3.1528	3.169	3.2123
4.3562	4.3928	4.4088	4.399	4.3897	4.3806	4.3884	4.4036	4.4135	4.4346	4.4789

Future Work

- Evaluate RNN for time-series analysis.
- Integrate additional attributes regarding the accident location (the count of stop signals and traffic lights, whether it is a parking lot or not ,etc...)
- Train Random Forest with smaller geohash zones, more specific predictions.